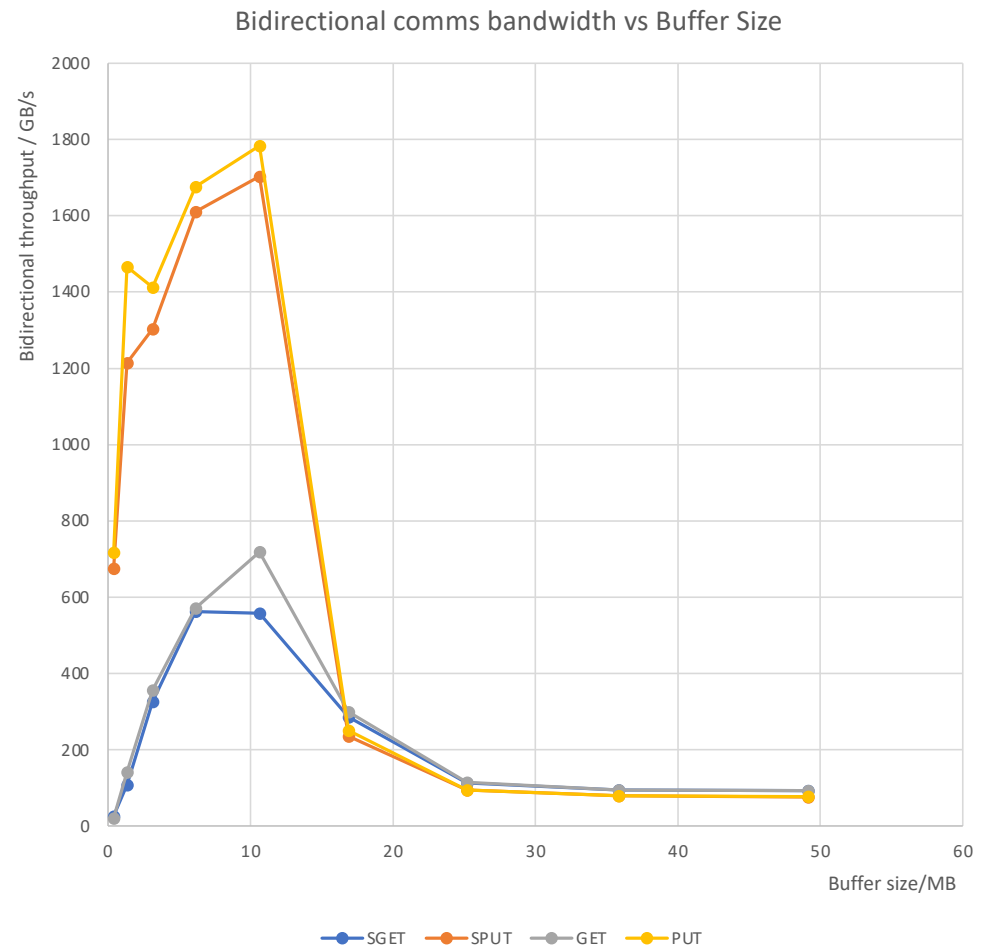


Grid performance summary: Benchmark_dwf

- Day 1: Baseline performance results
 - Machines configured in NPS4
 - Best results results obtained with
 - One MPI process per NUMA domain (1 quadrant/2CCDs)
 - Each MPI process running 16 threads (one per CPU)
 - 2.86 TFLOPS (compute-only)
 - 1.1 TFLOPS (including comms)
 - Single-node, $16^3 \times 32$ Lattice
 - Evenly distributed MPI (2 x 2 x 2 x 2)
- Day 1: observations
 - Diagnosing where bandwidth losses occur non-trivial
 - AOCC21 vs AOCC22 minimal improvement (~2%?)

Day 2 - Infinity

- Rome ∞ 32/16 bit Read/Write
 - Naples ∞ /16 bit Read/Write
- Expectations:
 - We expect to saturate DRAM
 - ... but we don't
 - 90GB/s we see here equivalent to 180GB/s on DRAM, i.e. approx. 50% DRAM peak
 - We expect GET to be faster than PUT
 - But it's not (when we hit cache)
 - GET=read remote, write local
- Findings
 - Streaming intrinsics marginally worse
 - Up to L3 cache size
 - High throughput
 - (S)PUT many x faster than (S)GET (why ?)
 - Above cache size
 - (S)GET 10%? faster than (S)PUT

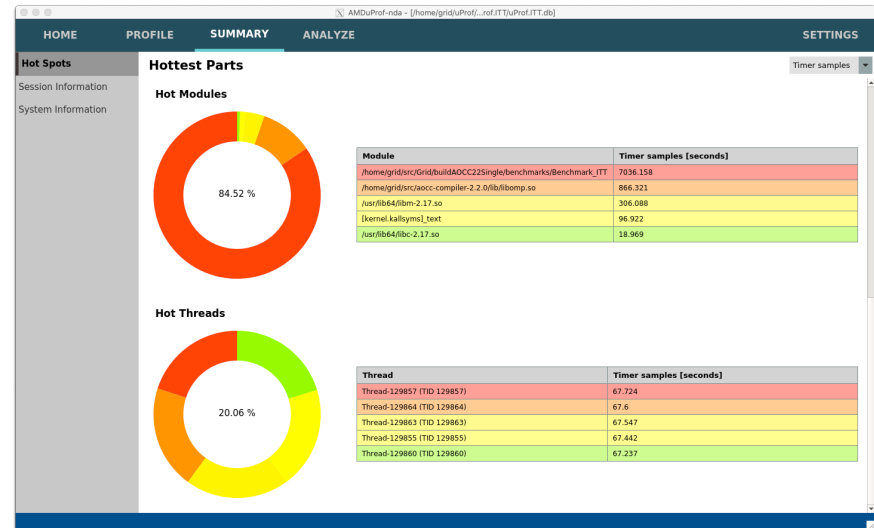
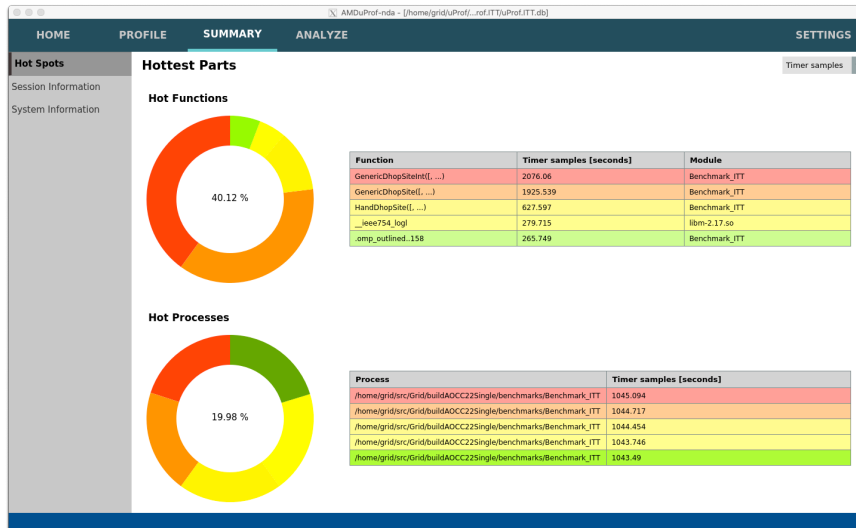


Day 2 - Other

- NPS 4 tests continued ...
- Compiler flags `-march=znver2`
 - 16 x 16 x 32 x 32 Lattice
 - 1 x 2 x 2 x 2 MPI
 - 3.1 TFLOPS (compute-only)
 - 1.2 TFLOPS (including comms)
- Optimising MPI Layout
 - 16 x 16 x 16 x 64 Lattice
 - 1 x 1 x 1 x 8 MPI... cheating
 - 2.4 TFLOPS (compute-only)
 - 1.5 TFLOPS (including comms)
- NPS1
 - 2 MPI processes 64 threads/proc
 - 16 x 16 x 16 x 64 Lattice
 - MPI layout 1 x 1 x 1 x 2 Lattice
 - 2.4 TFLOPS (compute-only)
 - 1.9 TFLOPS (including comms)
 - Needed:
 - `KMP_AFFINITY=verbose,granularity=fine,compact`
- Observations
 - NPS1 looks to be optimal
 - Because we avoid interior comms
 - ... i.e. buffer copies
 - Infinity link working
 - Might see 2 x boost of interior comms?
 - Cache-resident performance excellent

Final Summary::Highest performance obtained under NPS1::GFLOP/s rating of=1900

Appendix - AMDuProf



- We tried various tools, including the profiler
- We appreciated the technical support during the Hackathon